


Generating balanced multiplayer game levels with GANs and RL

Grzegorz Opoka (grzego.opoka@gmail.com)



AGH

Problem

Given players with different abilities, can we generate game levels that do not favor any of them?

Idea

Use RL agents to evaluate *fairness* (i.e. how balanced the level is) of given game level and then train GANs to generate game levels that optimize for it.

What's fairness?

Assuming all players are equally skilled. We say that a level is fair if it is not biased toward any player. We estimate *fairness* with following equation:

$$\text{fairness}(b) = - \sum_{a \in A} \left(\underbrace{\left(\frac{1}{N} \sum \text{game}(b, a) \right)}_{\text{games won by player } a \in A \text{ on average}} - \underbrace{\frac{1}{|A|}}_{\text{fair win-rate of game level}} \right)$$

The closer the value of *fairness* is to 0 the fairer the level is.


How to optimize fairness?

- Assuming RL agents are trained properly, they can be used as a replacement for human players. We can therefore use them to estimate *fairness*.
- We then simultaneously train Discriminator to predict the *fairness* estimated by Agents, and use it's knowledge about what makes levels fair to train Generator.

Agents

- Serve as replacement for slow evaluation by human players
- Used to estimate *fairness* of given game level
- Trained by PPO (Proximal Policy Optimization) algorithm to drive as best as possible

Different abilities (i.e. cars in our game)




Environment (Game)


- Basic race game where each players is starting from the same position and has to finish the race as fast as possible.

Actions

- Any combination of 4 actions (+ noop)




Rewards



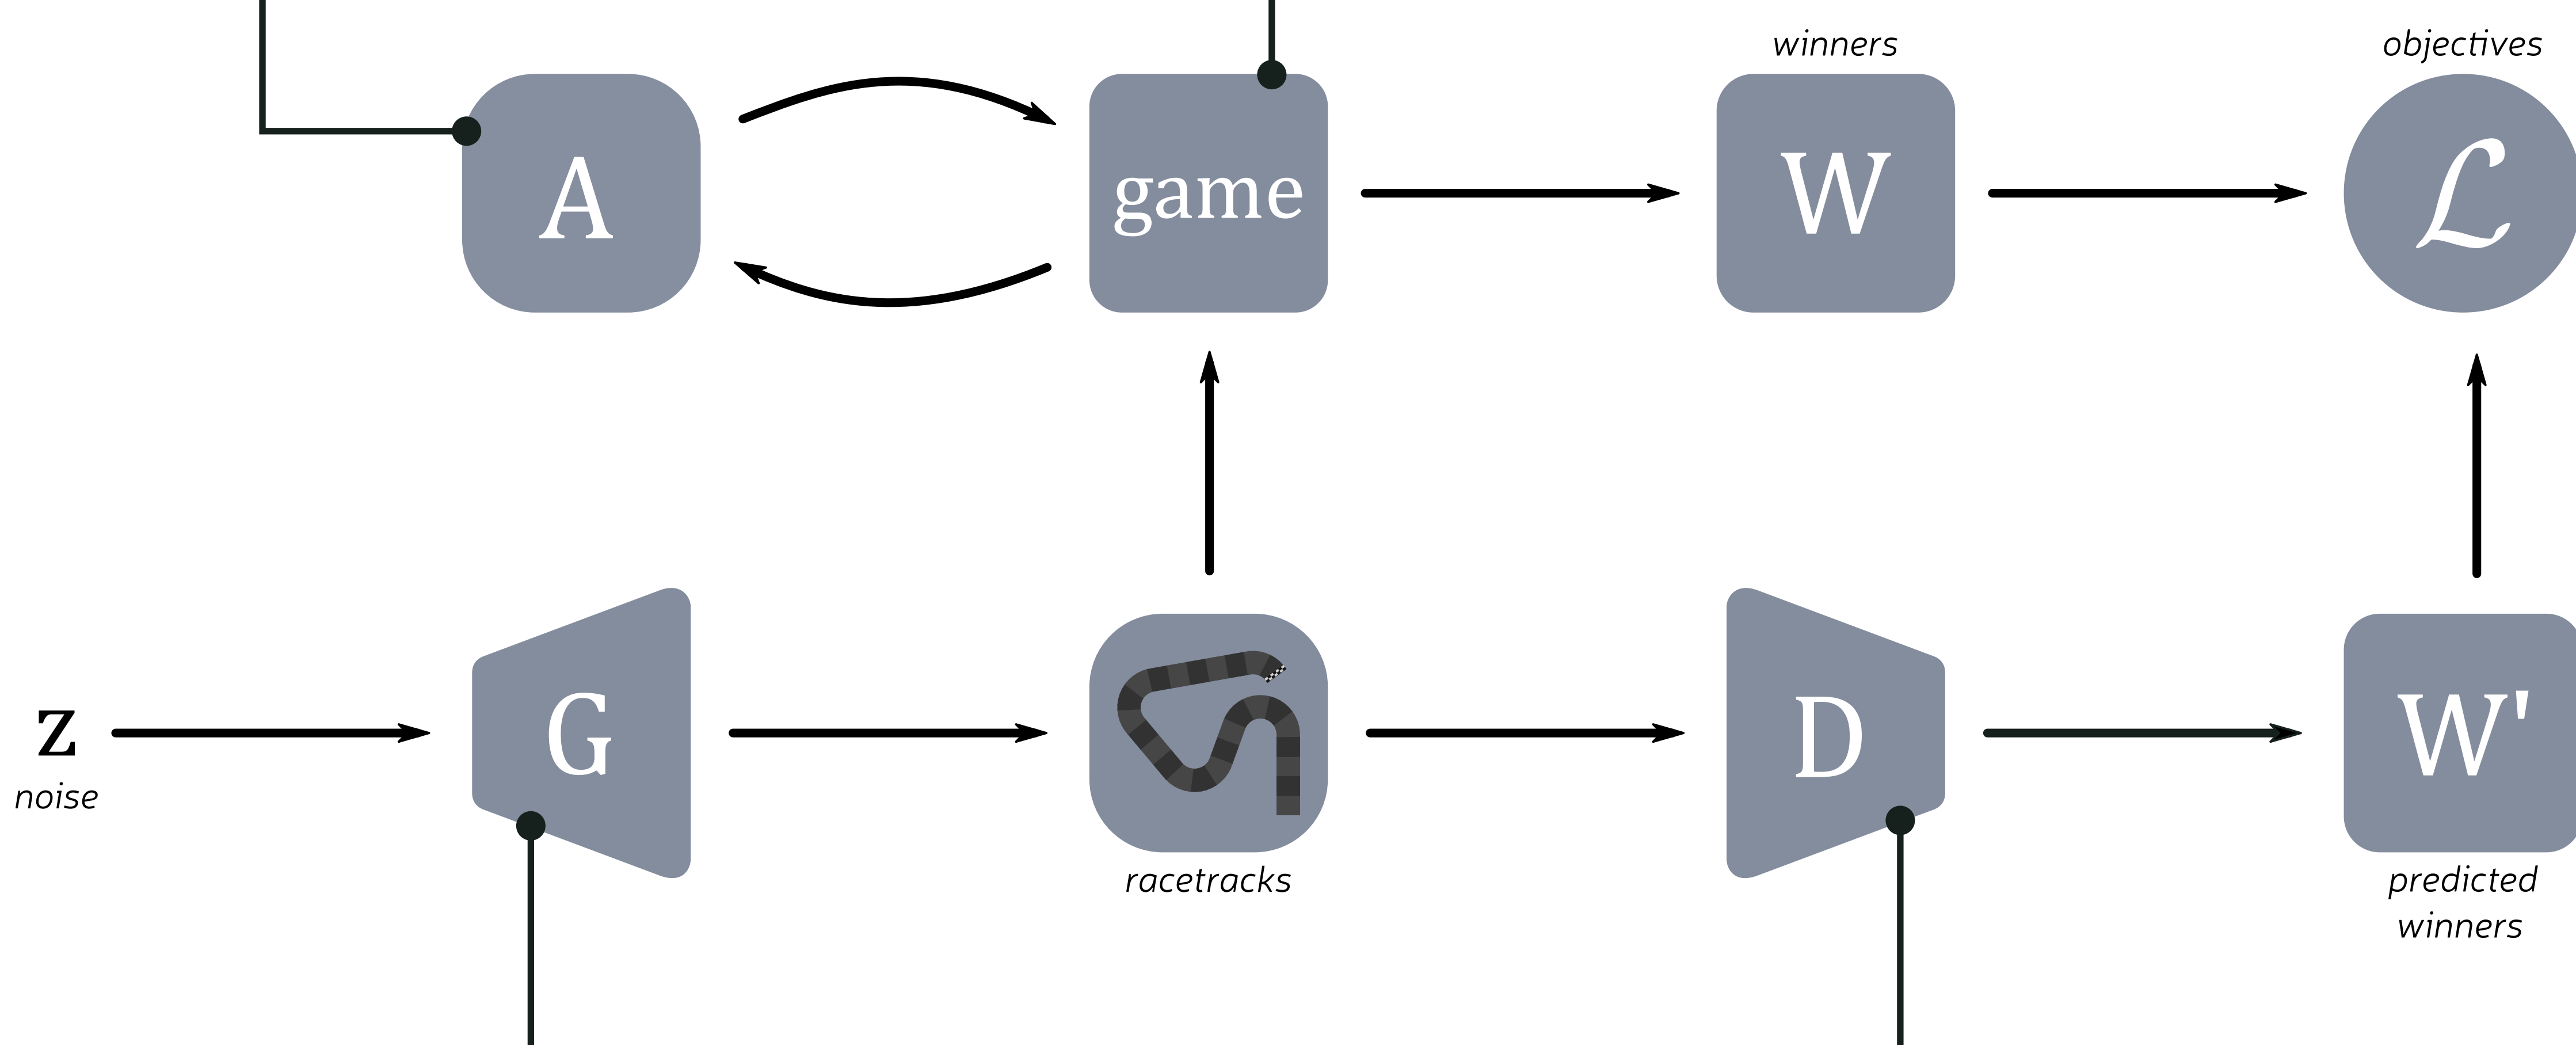

Racetracks

- Are represented as vectors of numbers. Each number corresponds to an angle at given point on track.




Observation (Input)

- Distances in 19 different directions.



Generator




Objective

- Generator goal is to maximize *fairness* of generated levels by minimizing following loss:

$$\mathcal{L}^G(\theta) = \mathbb{E}_{\mathbf{z} \sim \mathcal{N}} \left[-\frac{1}{|A|} \cdot \log(D(G_\theta(\mathbf{z}))) \right]$$

Discriminator



Objective

- Discriminator goal is to predict winner of given game level and by that approximate *fairness*. This is achieved by minimizing:

$$\mathcal{L}^D(\theta) = \mathbb{E}_{\mathbf{z} \sim \mathcal{N}} [-\text{game}(b, A) \cdot \log(D_\theta(G(\mathbf{z})))]$$

Algorithm

- Initialize all networks, A, G, D
- while** agents not trained **do**
- Reset environment
- while** episode not finished **do**
- Perform actions decided by agents, A
- Gather rewards received by agents
- Update state of environment
- end while**
- Use PPO to update agents, A , policies with data gathered during episode.
- end while**
- for** number of training iterations **do**
- Sample minibatch of noise vectors and generate boards, $\mathbf{z} \sim \mathcal{N}, b = G_\theta(\mathbf{z})$
- Simulate game and obtain winners, $w = \text{game}(b, A)$
- Update the discriminator, D_θ , parameters by descending gradient of:

$$\nabla \mathbb{E}_b [-w \cdot \log(D_\theta(b))]$$

- Sample noise vectors, $\mathbf{z} \sim \mathcal{N}$
- Update the generator, G_θ , parameters by descending gradient of:

$$\nabla \mathbb{E}_z \left[-\frac{1}{|A|} \cdot \log(D(G_\theta(\mathbf{z}))) \right]$$

17: **end for**

Further work

- Due to catastrophic forgetting Agents are not trained simultaneously with Generator & Discriminator, maybe there is a way to avoid it and train them together.
- Find other ways to include Agents feedback in generation process.
- Investigate possible extensions to generator loss function to encourage more diverse results.

Evaluation & Results

